

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
import keras
from keras.models import Sequential
from keras.layers import Dense, Conv2D, MaxPool2D, Flatten, Dropout
from keras.preprocessing.image import ImageDataGenerator
from keras.optimizers import adam_v2
```

```
from sklearn.metrics import classification_report, confusion_matrix
```

```
import tensorflow as tf
```

```
import cv2
import os
```

```
import numpy as np
```

```
labels = ['OSTRICH', 'PEACOCK']
img_size = 224
```

```
def get_data(data_dir):
```

```
    data = []
```

```
    for label in labels:
```

```
        path = os.path.join(data_dir, label)
```

```
        class_num = labels.index(label)
```

```
        print(path)
```

```
        for img in os.listdir(path):
```

```
            try:
```

```
                img_arr = cv2.imread(os.path.join(path, img))[...,:-1] #convert BGR to RGB f
```

```
                print(img_arr)
```

```
                resized_arr = cv2.resize(img_arr, (img_size, img_size)) # Reshaping images to
```

```
                data.append([resized_arr, class_num])
```

```
            except Exception as e:
```

```
                print(e)
```

```
    return np.array(data)
```

```
#Now we can easily fetch our train and validation data.
```

```
train = get_data('/content/drive/MyDrive/input/train')
```

```
val = get_data('/content/drive/MyDrive/input/test')
```

```
[[147 161 40]
 [150 164 43]
 [159 173 50]
 ...
 [ 43 152 73]
 [ 31 128 57]
 [ 52 144 77]]]
[[[160 151 134]
 [130 127 108]
 [100 105 83]
 ...
 [ 84 95 53]
 [109 112 83]
 [ 83 83 59]]

[[100 93 77]
 [120 117 98]
 [137 142 120]
 ...
 [115 123 82]
 [100 103 72]
 [105 108 81]]

[[102 100 85]
 [ 82 83 65]
 [118 123 101]
 ...
 [114 121 79]
 [ 85 91 55]
 [ 85 91 57]]

...

[[117 141 57]
 [104 129 45]
 [108 131 49]
 ...
 [125 148 76]
 [ 95 119 45]
 [ 97 121 47]]

[[ 91 117 30]
 [101 129 44]
 [107 132 49]
 ...
 [118 140 67]
 [ 90 112 37]
 [110 133 55]]

[[ 94 122 35]
 [ 90 118 33]
 [ 84 112 28]
 ...
 [ 81 103 30]
```

```
[ 96 119  41]
[107 130  50]]]
```

```
x_train = []
y_train = []
x_val = []
y_val = []

for feature, label in train:
    x_train.append(feature)
    y_train.append(label)

for feature, label in val:
    x_val.append(feature)
    y_val.append(label)

# Normalize the data
x_train = np.array(x_train) / 255
x_val = np.array(x_val) / 255

x_train.reshape(-1, img_size, img_size, 1)
y_train = np.array(y_train)

x_val.reshape(-1, img_size, img_size, 1)
y_val = np.array(y_val)

datagen = ImageDataGenerator(
    featurewise_center=False, # set input mean to 0 over the dataset
    samplewise_center=False, # set each sample mean to 0
    featurewise_std_normalization=False, # divide inputs by std of the dataset
    samplewise_std_normalization=False, # divide each input by its std
    zca_whitening=False, # apply ZCA whitening
    rotation_range = 30, # randomly rotate images in the range (degrees, 0 to 180)
    zoom_range = 0.2, # Randomly zoom image
    width_shift_range=0.1, # randomly shift images horizontally (fraction of total width)
    height_shift_range=0.1, # randomly shift images vertically (fraction of total height)
    horizontal_flip = True, # randomly flip images
    vertical_flip=False) # randomly flip images

datagen.fit(x_train)

model = Sequential()
model.add(Conv2D(32,3,padding="same", activation="relu", input_shape=(224,224,3)))
model.add(MaxPool2D())

model.add(Conv2D(32, 3, padding="same", activation="relu"))
model.add(MaxPool2D())
```

```
model.add(Conv2D(64, 3, padding="same", activation="relu"))
model.add(MaxPool2D())
model.add(Dropout(0.4))

model.add(Flatten())
model.add(Dense(128,activation="relu"))
model.add(Dense(2, activation="softmax"))

model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 224, 224, 32)	896
max_pooling2d (MaxPooling2D)	(None, 112, 112, 32)	0
conv2d_1 (Conv2D)	(None, 112, 112, 32)	9248
max_pooling2d_1 (MaxPooling2D)	(None, 56, 56, 32)	0
conv2d_2 (Conv2D)	(None, 56, 56, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 28, 28, 64)	0
dropout (Dropout)	(None, 28, 28, 64)	0
flatten (Flatten)	(None, 50176)	0
dense (Dense)	(None, 128)	6422656
dense_1 (Dense)	(None, 2)	258

Total params: 6,451,554
Trainable params: 6,451,554
Non-trainable params: 0

```
from keras.optimizer_v2.adam import Adam
opt = Adam(learning_rate=0.000001)
model.compile(optimizer = opt , loss = tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True))

history = model.fit(x_train,y_train,epochs = 150 , validation_data = (x_val, y_val))
```

```
10/10 [=====] - 20s 2s/step - loss: 0.2585 - accuracy: 0.908
Epoch 124/150
10/10 [=====] - 20s 2s/step - loss: 0.2541 - accuracy: 0.924
Epoch 125/150
10/10 [=====] - 20s 2s/step - loss: 0.2529 - accuracy: 0.921
Epoch 126/150
10/10 [=====] - 20s 2s/step - loss: 0.2527 - accuracy: 0.927
Epoch 127/150
10/10 [=====] - 20s 2s/step - loss: 0.2482 - accuracy: 0.905
Epoch 128/150
10/10 [=====] - 20s 2s/step - loss: 0.2497 - accuracy: 0.918
Epoch 129/150
10/10 [=====] - 20s 2s/step - loss: 0.2460 - accuracy: 0.921
Epoch 130/150
10/10 [=====] - 20s 2s/step - loss: 0.2477 - accuracy: 0.911
Epoch 131/150
10/10 [=====] - 20s 2s/step - loss: 0.2475 - accuracy: 0.921
Epoch 132/150
10/10 [=====] - 20s 2s/step - loss: 0.2453 - accuracy: 0.921
Epoch 133/150
10/10 [=====] - 20s 2s/step - loss: 0.2420 - accuracy: 0.924
Epoch 134/150
10/10 [=====] - 20s 2s/step - loss: 0.2421 - accuracy: 0.914
Epoch 135/150
10/10 [=====] - 20s 2s/step - loss: 0.2349 - accuracy: 0.927
Epoch 136/150
10/10 [=====] - 20s 2s/step - loss: 0.2350 - accuracy: 0.918
Epoch 137/150
10/10 [=====] - 20s 2s/step - loss: 0.2417 - accuracy: 0.924
Epoch 138/150
10/10 [=====] - 20s 2s/step - loss: 0.2320 - accuracy: 0.933
Epoch 139/150
10/10 [=====] - 20s 2s/step - loss: 0.2365 - accuracy: 0.918
Epoch 140/150
10/10 [=====] - 20s 2s/step - loss: 0.2325 - accuracy: 0.921
Epoch 141/150
10/10 [=====] - 20s 2s/step - loss: 0.2296 - accuracy: 0.924
Epoch 142/150
10/10 [=====] - 20s 2s/step - loss: 0.2325 - accuracy: 0.933
Epoch 143/150
10/10 [=====] - 20s 2s/step - loss: 0.2336 - accuracy: 0.918
Epoch 144/150
10/10 [=====] - 20s 2s/step - loss: 0.2274 - accuracy: 0.918
Epoch 145/150
10/10 [=====] - 20s 2s/step - loss: 0.2264 - accuracy: 0.930
Epoch 146/150
10/10 [=====] - 20s 2s/step - loss: 0.2290 - accuracy: 0.921
Epoch 147/150
10/10 [=====] - 20s 2s/step - loss: 0.2277 - accuracy: 0.924
Epoch 148/150
10/10 [=====] - 20s 2s/step - loss: 0.2247 - accuracy: 0.921
Epoch 149/150
10/10 [=====] - 20s 2s/step - loss: 0.2285 - accuracy: 0.921
Epoch 150/150
10/10 [=====] - 20s 2s/step - loss: 0.2212 - accuracy: 0.924
```

```



acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']

epochs_range = range(150)

plt.figure(figsize=(15, 15))
plt.subplot(2, 2, 1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')

plt.subplot(2, 2, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()

```



```

predictions = np.argmax(model.predict(x_val), axis=-1),
predictions = np.asarray(predictions)
predictions = predictions.reshape(1,-1)[0]
print(classification_report(y_val, predictions, target_names = ['Ostrich (Class 0)', 'Peacock

```

	precision	recall	f1-score	support
Ostrich (Class 0)	1.00	0.80	0.89	5
Peacock (Class 1)	0.83	1.00	0.91	5
accuracy			0.90	10
macro avg	0.92	0.90	0.90	10
weighted avg	0.92	0.90	0.90	10

Task 1: Run the above code with given dataset.

Task 2: Run the code with different dataset

[Colab paid products](#) - [Cancel contracts here](#)

✓ 0s completed at 01:01

